CS 124/LINGUIST 180
From Languages to Information

Unix for Poets

Dan Jurafsky

(original by Ken Church, modifications by me and Chris Manning)

Stanford University

Unix for Poets

- Text is everywhere
  - The Web
  - Dictionaries, corpora, email, etc.
  - Billions and billions of words
- What can we do with it all?
- It is better to do something simple, than nothing at all.
- You can do simple things from a Unix command-line
- Sometimes it's much faster even than writing a quick python tool
- DIY is very satisfying

Exercises we'll be doing today

1. Count words in a text
2. Sort a list of words in various ways
   - ascii order
   - ''rhyming'' order
3. Extract useful info from a dictionary
4. Compute ngram statistics

Tools

- **grep**: search for a pattern (regular expression)
- **sort**
- **uniq –c** (count duplicates)
- **tr** (translate characters)
- **wc** (word – or line – count)
- **sed** (edit string -- replacement)
- **cat** (send file(s) in stream)
- **echo** (send text in stream)

- **cut** (columns in tab-separated files)
- **paste** (paste columns)
- **head**
- **tail**
- **rev** (reverse lines)
- **comm**
- **join**

4

Prereq: If you are on a Mac:

- Open the Terminal app

Prereq: If you are on a Windows 10 machine and don't have Ubuntu on your machine:

- For today's class, it's easiest to work with someone who has a Mac or Linux machine, or has Ubuntu already.

- Or you can do the following so that you will have this ability:

- Watch the first 9 minutes of Bryan's lovely pa0 video about how to download and install Ubuntu:
  https://canvas.stanford.edu/courses/144170/modules/items/981067

- Watch Chris Gregg's excellent UNIX videos here: Logging in, the first 7 "File System" videos, and the first 8 "useful commands" videos.
  https://web.stanford.edu/class/archive/cs/cs107/cs107.1186/unixref/

- From there you can use the ssh command to connect to the myth machines. Just be sure to keep track in your own mind of whether you're on myth or your own laptop at any given moment! The ssh command you want to type is:

- ssh [sunet]@rice.stanford.edu where [sunet] is your SUNet ID. It will ask for your password, which is your usual Stanford password, and you will have to do two-step authentication.

6

Prerequisites: get the text file we are using

- rice: ssh into a rice or myth and then do (don't forget the final ".")

```
cp /afs/ir/class/cs124/WWW/nyt_200811.txt .
```

- Or download to your own Mac or Unix laptop this file:

http://cs124.stanford.edu/nyt_200811.txt

```
Or:
```

```
scp cardinal:/afs/ir/class/cs124/WWW/nyt_200811.txt  .
```

Prerequisites

- The unix "man" command
  - e.g., `man tr`
  - Man shows you the command options; it's not particularly friendly

8

Prerequisites

- How to chain shell commands and deal with input/output
- Input/output redirection:
  - `>` "output to a file"
  - `<` "input from a file"
  - `|` "pipe"
- CTRL-C

9

OK, you're ready to start!

Exercise 1: Count words in a text

- Input: text file (nyt_200811.txt)

- Output: list of words in the file with freq counts

- Algorithm
  1. Tokenize (`tr`)
  2. Sort (`sort`)
  3. Count duplicates (`uniq —c`)

- Go read the man pages and figure out how to pipe these together

Solution to Exercise 1

- ```
  tr -sc 'A-Za-z' '\n' < nyt_200811.txt |
  sort | uniq -c
  ```

```
633 A
  1 AA
  1 AARP
  1 ABBY
 41 ABC
  1 ABCNews
```

(Do you get a different sort order? In some versions of UNIX, sort doesn't use ASCII order (uppercase before lowercase).)

Some of the output

- ```
  tr -sc 'A-Za-z' '\n'
  < nyt_200811.txt |
  sort | uniq -c |
  head –n 5
  ```

```
633 A
  1 AA
  1 AARP
  1 ABBY
 41 ABC
```

- ```
  tr -sc 'A-Za-z' '\n'
  < nyt_200811.txt |
  sort | uniq -c |
  head
  ```

- **head** gives you the first 10 lines

- **tail** does the same with the end of the input

- (You can omit the "-n" but it's discouraged.)

Extended Counting Exercises

1. Merge upper and lower case by downcasing everything
    - Hint: Put in a second tr command

2. How common are different sequences of vowels (e.g., the sequences "ieu" or just "e" in "lieutenant")?
    - Hint: Put in a second tr command

14

Solutions

Merge upper and lower case by downcasing everything

```
tr -sc 'A-Za-z' '\n' < nyt_200811.txt | tr 'A-Z' 'a-z'
| sort | uniq -c
or
tr -sc 'A-Za-z' '\n' < nyt_200811.txt | tr '[:upper:]'
'[:lower:]' | sort | uniq -c
```

1. tokenize by replacing the complement of letters with newlines
2. replace all uppercase with lowercase
3. sort alphabetically
4. merge duplicates and show counts

Solutions

- How common are different sequences of vowels (e.g., ieu)
- `tr 'A-Z' 'a-z' < nyt_200811.txt | tr -sc 'aeiou' '\n' | sort | uniq -c`

Sorting and reversing lines of text

- `sort`
- `sort —f`   Ignore case
- `sort —n`   Numeric order
- `sort —r`   Reverse sort
- `sort —nr`   Reverse numeric sort

- `echo "Hello" | rev`

17

Counting and sorting exercises

- Find the 50 most common words in the NYT
  - Hint: Use sort a second time, then head

- Find the words in the NYT that end in "zz"
  - Hint: Look at the end of a list of reversed words
  - tr 'A-Z' 'a-z' < filename | tr –sc 'a-z' '\n' | rev | sort | rev | uniq -c

Counting and sorting exercises

- Find the 50 most common words in the NYT

```
tr -sc 'A-Za-z' '\n' < nyt_200811.txt |
sort | uniq -c | sort -nr | head -n 50
```

- Find the words in the NYT that end in "zz"

```
tr -sc 'A-Za-z' '\n' < nyt_200811.txt | tr
'A-Z' 'a-z' | rev | sort | uniq -c | rev |
tail -n 10
```

Lesson

- Piping commands together can be simple yet powerful in Unix
- It gives flexibility.

- Traditional Unix philosophy: small tools that can be composed

20

Bigrams = word pairs and their counts

Algorithm:

1. Tokenize by word

2. Create two almost-duplicate files of words, off by one line, using **tail**

3. **paste** them together so as to get $word_i$ and $word_{i+1}$ on the same line

4. Count

Bigrams

- `tr -sc 'A-Za-z' '\n' < nyt_200811.txt > nyt.words`
- `tail -n +2 nyt.words > nyt.nextwords`
- `paste nyt.words nyt.nextwords > nyt.bigrams`
- `head —n 5 nyt.bigrams`

```
KBR     said
said    Friday
Friday  the
the     global
global  economic
```

22

Exercises

- Find the 10 most common bigrams
  - (For you to look at:) What part-of-speech pattern are most of them?

- Find the 10 most common trigrams

Solutions

- Find the 10 most common bigrams

```
tr 'A-Z' 'a-z' < nyt.bigrams | sort | uniq
-c | sort -nr | head -n 10
```

- Find the 10 most common trigrams

```
tail -n +3 nyt.words > nyt.thirdwords
```

```
paste nyt.words nyt.nextwords nyt.thirdwords >
nyt.trigrams
```

```
cat nyt.trigrams | tr "[:upper:]" "[:lower:]" | sort |
uniq -c | sort -rn | head -n 10
```

grep

- Grep finds patterns specified as regular expressions
- `grep rebuilt nyt_200811.txt`

Conn and Johnson, has been rebuilt, among the first of the 222

move into their rebuilt home, sleeping under the same roof for the

the part of town that was wiped away and is being rebuilt. That is

to laser trace what was there and rebuilt it with accuracy," she

home - is expected to be rebuilt by spring. Braasch promises that a

25

grep

- Grep finds patterns specified as regular expressions
  - **g**lobally search for **r**egular **e**xpression and **p**rint

- Finding words ending in –ing:
- `grep 'ing$' nyt.words |sort | uniq —c`

grep

- grep is a filter – you keep only some lines of the input
- `grep gh`   keep lines containing ''gh''
- `grep 'ˆcon'`   keep lines beginning with ''con''
- `grep 'ing$'`   keep lines ending with ''ing''
- `grep —v gh`   keep lines NOT containing "gh"

grep versus egrep (grep –E)

- `egrep or grep –E` [extended syntax]
- In egrep, +, ?, |, (, and) are automatically metacharacters
- In grep, you have to backslash them
- To find words ALL IN UPPERCASE:
- `egrep '^[A-Z]+$' nyt.words |sort|uniq -c`
- `== grep '^[A-Z]\+$' nyt.words |sort|uniq -c`

`(confusingly on some systems grep acts like egrep)`

Counting lines, words, characters

- ```
  wc nyt_200811.txt
  ```
  ```
    70334   509851 3052306 nyt_200811.txt
  ```


- ```
  wc -l nyt.words
  ```
  ```
    70334 nyt_200811.txt
  ```

Exercise: Why is the number of words different?

Exercises on grep & wc

- How many all uppercase words are there in this NYT file?

- How many 4-letter words?

- How many different words are there with no vowels
  - What subtypes do they belong to?

- How many "1 syllable" words are there
  - That is, ones with exactly one sequence of vowels

Type/token distinction: different words (types) vs. instances (tokens)

Solutions on grep & wc

- How many all uppercase words are there in this NYT file?

```
grep -E '^[A-Z]+$' nyt.words | wc
```

- How many 4-letter words?

```
grep -E '^[a-zA-Z]{4}$' nyt.words | wc
```

- How many different words are there with no vowels

```
grep -v '[AEIOUaeiou]' nyt.words | sort | uniq | wc
```

- How many "1 syllable" words are there

- ```
  tr 'A-Z' 'a-z' < nyt.words | grep -P
  '^[^aeiou]*[aeiou]+[^aeiou]*$' | uniq | wc
  ```

Type/token distinction: different words (types) vs. instances (tokens)

sed

- sed is used when you need to make systematic changes to strings in a file (larger changes than 'tr')

- It's line based: you optionally specify a line (by regex or line numbers) and specific a regex substitution to make

- For example to change all cases of "George" to "Jane":


- ```
  sed 's/George/Jane/' nyt_200811.txt | less
  ```

32

sed exercises

- Count frequency of word initial consonant sequences
  - Take tokenized words
  - Delete the first vowel through the end of the word
  - Sort and count

- Count word final consonant sequences

sed exercises

- Count frequency of word initial consonant sequences

```
tr "[:upper:]" "[:lower:]" < nyt.words | sed
's/[aeiou].*$//' | sort | uniq -c
```

- Count word final consonant sequences

```
tr "[:upper:]" "[:lower:]" < nyt.words | sed
's/^.*[aeiou]//' | sort | uniq -c | sort -rn
| less
```

34

Extra Credit – Secret Message

- Now, let's get some more practice with Unix!

- The answers to the extra credit exercises will reveal a secret message.

- We will be working with the following text file for these exercises: https://cs124.stanford.edu/secret_ec.txt

- To receive credit, enter the secret message here: https://forms.gle/57okKzZzWeijP4RL7

Extra Credit Exercise 1

- Find the 2 most common words in secret_ec.txt containing the letter e.

- Your answer will correspond to the first two words of the secret message.

Extra Credit Exercise 2

- Find the 2 most common bigrams in secret_ec.txt where the second word in the bigram ends with a consonant.

- Your answer will correspond to the next four words of the secret message.

37

Extra Credit Exercise 3

- Find all 5 letter words that only appear once in secret_ec.txt.

- Concatenate your result. This will be the final word of the secret message.